

# Origin Axiom — Phase 3 (Exploratory Add-on): Flavor Phase Integration

Origin Axiom Collaboration

## Abstract

Phase 3 integrates an empirically fitted phase parameter  $\theta$  (from CKM+PMNS within a declared ansatz) into the Phase 2 vacuum residue mechanism, and documents the resulting correlation under strict under-claiming.

## 1 Scope and Non-Claims

### 1.1 Role of Phase 3 in the program

Phase 3 is a *ledger-compatible add-on*: it extracts a candidate value and uncertainty interval for  $\theta$  by fitting a fixed, explicitly stated ansatz to frozen external flavor-sector targets (CKM and PMNS CP-phase information). The Phase 3 deliverable is not the internal fit machinery; it is the exported, schema-stable artifact `phase_03_theta_filter.json`, designed to be consumable by the Phase 0 corridor/ledger aggregation without requiring Phase 3 internals.

Concretely:

- **Inputs:** external target snapshot encoded in `phase3/fit/targets.yaml` (treated as frozen for this Phase 3 run) and the fixed ansatz implementation.
- **Outputs:** (i) fit summary and diagnostics tables, (ii) figures, (iii) this PDF report, and (iv) the ledger artifact `phase_03_theta_filter.json`.

### 1.2 What Phase 3 does and does not establish

Phase 3 establishes only the following narrow fact: under the Phase 3 baseline ansatz and frozen targets, there exists a best-fit  $\theta$  and an uncertainty interval produced by an explicit, reproducible scan.

**Non-claims (explicit).** Phase 3 does *not* claim:

- a derivation of a unique fundamental constant  $\theta_*$ ;
- a proof that the Origin Axiom is correct, complete, or uniquely determined;
- a reduction of Standard Model free parameters, or a predictive theory of the CKM/PMNS structures;
- a measurement-level global fit (this is not a replacement for PDG/NuFIT analyses);
- robustness to arbitrary re-parameterizations or model classes beyond the declared ansatz family.

### 1.3 Success criteria (Phase 1/2 style contract)

A Phase 3 run is considered successful iff:

- the workflow gate reproduces all declared artifacts end-to-end;
- the bundle manifest lists exactly the declared inputs/outputs with hashes and sizes;
- the paper claims remain within the narrow scope above (no implied derivation or discovery).

**Ledger-level outcome.** At the level of the Phase 0 corridor ledger, Phase 3 contributes only a filter: it may further narrow the corridor defined by earlier phases or, in some configurations, eliminate it entirely. In the baseline configuration documented here (ansatz locked in `phase3/fit/targets.yaml` with the stated CKM/PMNS snapshot), the Phase 0 ledger reports that the combined corridor is empty once the Phase 3 filter is applied. We interpret this as a negative result for that specific ansatz/target combination, not as a proof that no Origin-Axiom-compatible corridor exists.

## 2 Fit Pipeline: CKM+PMNS $\rightarrow \theta$

We define an explicit mapping (ansatz) from a single phase parameter  $\theta$  to a set of target observables and perform a grid-based  $\chi^2$  scan. We report the best-fit  $\theta$ , an uncertainty interval, and diagnostics. External CKM/PMNS

targets are treated as frozen inputs for this Phase 3 add-on, sourced from standard references (PDG and NuFIT snapshots) and encoded in `phase3/fit/targets.yaml` [1, 2].

## 2.1 Objective function

Let the ansatz predict a set of phase-like targets  $\{\varphi_i(\theta)\}$  to be compared against frozen targets  $\{\varphi_i^{\text{tgt}}\}$  with uncertainties  $\{\sigma_i\}$ . Because phases are defined modulo  $2\pi$ , we use a wrapped distance

$$\Delta(\alpha, \beta) = \text{wrap}(\alpha - \beta) \in (-\pi, \pi], \quad \chi^2(\theta) = \sum_i \left( \frac{\Delta(\varphi_i(\theta), \varphi_i^{\text{tgt}})}{\sigma_i} \right)^2.$$

The implementation details (grid bounds, step size, and target keys) are recorded in the accompanying metadata JSON produced by the workflow.

## 2.2 Discrete offset sweep and baseline choice

To avoid hidden parameter fitting, the Phase 3 baseline fixes all offset choices a priori and fits only  $\theta$ . We nevertheless test a small discrete set of candidate fixed offsets for the PMNS phase mapping,  $b_{\text{PMNS}} \in \{0, \pi, -\pi/2, +\pi/2\}$ , holding the rest of the ansatz fixed. For each hypothesis we run the same grid scan over  $\theta$  and compare the minimum  $\chi^2$  values. The results are reported in Appendix 1.

We adopt as the Phase 3 baseline the discrete hypothesis with the lowest minimum  $\chi^2$  under the frozen targets. For the current run this selects  $b_{\text{PMNS}} = \pi$ .

## 3 Injection Pipeline: $\theta \rightarrow$ Phase 2 diagnostic

The injection stage is a *downstream diagnostic* that consumes the fitted `theta_fit_summary.csv` artifact and evaluates the Phase 2 injection observable (vacuum-residue response) as a function of  $\theta$ . This stage does *not* participate in the Phase 3 fit objective and therefore cannot retroactively improve the CKM/PMNS fit. Its role is purely cross-phase: it allows the Phase 3 best-fit interval to be visualized against the Phase 2 diagnostic curve, producing Figure 2 and associated provenance metadata.

All injection outputs are packaged in the Phase 3 paper bundle and listed in the manifest. The workflow records the exact inputs, hashes, and run environment so the diagnostic can be regenerated verbatim.

## 4 Falsifiability and Failure Modes

Phase 3 is falsifiable in the narrow sense appropriate to an exploratory fit: the mapping from  $\theta$  to the frozen flavor targets can fail to produce a stable, interpretable interval.

### 4.1 Explicit failure conditions

We treat the Phase 3 extraction as *invalid* if any of the following occur:

- **Non-identifiability:** multiple disjoint  $\theta$  regions yield indistinguishably good minima (flat or multi-modal  $\chi^2$  without a defensible interval).
- **Offset ambiguity:** the discrete offset sweep yields competing hypotheses with comparable minima, preventing a pre-registered baseline choice from being meaningful.
- **Target drift collapse:** a reasonable update of the external targets (future PDG/NuFIT releases) removes the minimum or shifts it outside any stable corridor overlap with other phases.
- **Ansatz fragility:** small, clearly stated modifications to the mapping family (within the declared ansatz class) destroy the existence of a stable minimum, indicating the result is an artifact of parameterization rather than a durable extraction.
- **Ledger incompatibility:** when the Phase 3  $\theta$ -filter is injected into the Phase 0 corridor ledger, the resulting combined corridor becomes empty (no  $\theta$  survives simultaneous application of the Phase 0–2 and Phase 3 filters). This is treated as evidence against the locked Phase 3 ansatz/target combination until the ansatz or external targets are revised.

### 4.2 What Phase 3 can support downstream

If Phase 3 passes the checks above, it supports only this downstream use: the exported `phase_03_theta_filter.json` interval can be compared by Phase 0 ledger tooling to other phase constraints to test whether a corridor overlap emerges. No stronger inference is warranted.

## 5 Limitations

- **External-target dependence:** results are conditional on frozen PDG/NuFIT-style targets encoded in `phase3/fit/targets.yaml`. Updates to those targets may shift or erase the extracted interval.
- **Ansatz dependence:** the extraction is conditional on the declared mapping family. Phase 3 does not claim this family is unique or theoretically preferred.
- **Not a global fit:** this is not a full oscillation-data reanalysis or a replacement for PDG or NuFIT methodologies; it is a single-parameter exploratory scan.
- **Interpretation bound:** the only supported downstream interpretation is ledger comparison (corridor overlap) with other phases; no discovery claim is implied.

## A Claims Table (Phase 3)

See `phase3/CLAIMS_TABLE.md` for the live evidence map.

## B Reproducibility

Phase 3 follows the same gate structure as Phases 0–2, with explicit levels and a single Snakemake entry point. The authoritative, machine-readable reproducibility contract is documented in `phase3/REPRODUCIBILITY.md`, but we summarize the structure here.

### Levels.

- **Level A (repo snapshot):** verify the paper and bundle from an existing repository snapshot. This level checks that the committed artifacts and the bundle manifest are internally consistent. Entry point: `bash scripts/phase3_gate.sh --level A`.
- **Level B (regenerate artifacts):** rebuild all tables, figures, the paper, and the paper bundle from source using Snakemake, then verify the bundle. This is the default reproducibility target for external readers: `snakemake -s phase3/workflow/Snakefile -c 1 all`, followed by `bash scripts/phase3_gate.sh --level B`.

Table 1: Discrete fixed-offset sweep for  $b_{\text{PMNS}}$  with a single fitted parameter  $\theta$ . Lower  $\chi^2$  is better;  $\Delta\chi^2$  is relative to the best row.

$b_{\text{PMNS}}$ (deg)	$\theta^*$ (deg)	$\chi^2$	$\Delta\chi^2$	$\theta_{68\%}^{\text{lo}}$ (rad)	$\theta_{68\%}^{\text{hi}}$ (rad)
180	65.682	0.675474	0.000000	1.120292	1.172128
90	65.790	1.881895	1.206421	1.122491	1.174013
-90	65.556	9.093410	8.417936	1.118407	1.169929
0	65.916	12.712767	12.037293	1.124690	1.176212

- **Level C (heavy runs):** optional developer mode that can populate `phase3/outputs/runs/` (git-ignored) with additional scans or diagnostics. This mode is not required to reproduce any claims in this paper, and is intended only for extended exploration: `bash scripts/phase3_gate.sh --level C`.

**Bundle and evidence location.** All artifacts used as evidence for the Phase 3 claims are collected under `phase3/outputs/paper_bundle/`, with a `run_index.json` and `bundle_manifest.json` that enumerate inputs and outputs with hashes and sizes. The canonical PDF for this phase is `phase3/artifacts/origin-axiom-p` built by the Phase 3 Snakemake workflow.

## References

- [1] Particle Data Group. Review of particle physics. *Phys. Rev. D*, 110:030001, 2024.
- [2] I. Esteban, M. C. Gonzalez-Garcia, M. Maltoni, I. Martinez-Soler, and T. Schwetz. The fate of hints: updated global analysis of three-flavor neutrino oscillations. *JHEP*, 09:178, 2020. NuFIT 5.2 (online update; see [nu-fit.org](http://nu-fit.org)).

## C Phase 3 $\theta$ -filter artifact (Phase 0 ledger interface)

Phase 3 emits a machine-readable  $\theta$ -filter artifact: `phase3/outputs/theta_filter/phase_03_theta_filter.json`. This file is the Phase 3 contribution to the Phase 0 corridor method: it reports an admissible set of  $\theta$  values under an explicitly declared Phase 3 test suite, along with provenance sufficient for audit and reproduction.

## C.1 Declared test suite

Phase 3 is an empirical calibration filter (not an OA-binding demonstration). We therefore define a single test:

- **fit\_compat\_interval:**  $\theta$  is admissible if it lies within the declared fit interval reported in `theta_fit_summary.csv`.

This choice is intentionally conservative: it makes the corridor definition explicit and reproducible, and it avoids parameter proliferation or implicit re-optimization beyond the one-parameter scan.

## C.2 Schema compliance

The JSON conforms to the Phase 0 ledger interface: it declares a  $\theta$  domain on  $[0, 2\pi)$ , provides an interval-form corridor representation, and includes a deterministic grid+pass array representation, plus provenance bindings (commit/config/environment hashes when available).

## C.3 Ledger outcome in the current configuration

The Phase 0 ledger applies the Phase 3  $\theta$ -filter on top of the existing Phase 0–2 filters and records the resulting combined corridor in `phase0/phase_outputs/theta_corridor_history.jsonl`. In the baseline configuration documented here, the ledger reports that the combined corridor is empty once the Phase 3 filter is applied. This is logged as a negative test for the present Phase 3 ansatz/target combination, and it is intentionally encoded so that future runs with alternative ansatz choices or updated external targets can be compared against this outcome.