Phase 4: Vacuum-to-FRW Consistency and Scale Sanity

A corridor-style test of the Phase 3 global-amplitude mechanism

Origin Axiom Program

January 7, 2026

Abstract

Phase 4 tests whether the canonical global-amplitude mechanism defined in Phase 3 can be connected, in a structurally reasonable way, to FRW-style dynamics and vacuum-energy-like observables. The goal is to probe whether the non-cancellation floor that stabilises the toy vacuum can also support scale-sane behaviour when mapped into simple cosmological modules, without claiming a full cosmological fit or a theory of everything.

1 Introduction

Phase 4 sits between the mechanism-level vacuum work of Phase 3 and any future attempts to build a unified picture of vacuum, matter, and geometry. Its mission is narrow:

- take the canonical Phase 3 global-amplitude mechanism and non-cancellation floor as given;
- define simple, explicit mappings from the floor-enforced amplitude or residue into FRW-like dynamics and vacuum-energy-like observables; and
- study the resulting behaviour of the Origin-Axiom phase parameter θ in the corridor / ledger framework of Phase 0.

The guiding question is not whether we can reproduce the full ΛCDM model or fit precise cosmological parameters, but whether the Phase 3 mechanism can be made compatible with toy FRW modules in a way that is numerically stable, structurally sane, and expressible as a θ-filter in the sense of Phase 0.

Throughout this phase we distinguish carefully between:

- *binding* outputs, which may eventually define a Phase 4 θ-filter; and
- *non-binding* diagnostics and figures, which serve only as intuition and internal checks.

The present rung does not define any concrete mappings or filters. It only provides a minimal paper skeleton so that future rungs can add well-documented mechanisms and experiments without restructuring the front matter.

2 Mapping families: first pass (F1)

Phase 4 takes as input the Phase 3 global-amplitude mechanism: a toy vacuum with an unconstrained observable $A_0(\theta)$, a non-cancellation floor ε, and a floor-enforced amplitude $A(\theta) = \max(A_0(\theta), \varepsilon)$ defined on a grid $\theta \in [0, 2\pi)$. The present paper introduces a first, explicit mapping family, denoted **F1**, from this structure to a toy vacuum-energy-like scalar.

2.1 F1: direct scalar mapping from $A(\theta)$

The F1 family is intentionally simple. For a fixed Phase 3 vacuum configuration and floor ε (taken from the Phase 3 baseline diagnostics), we define a scalar

$$E_{\mathrm{vac}}(\theta) \;=\; \alpha\, A(\theta)^{\beta}, \tag{1}$$

where $\alpha > 0$ and $\beta > 0$ are explicit, configurable parameters. At this rung we adopt a conservative default, $\alpha = 1$ and $\beta = 2$, and focus on structural behaviour rather than numerical normalisation.

Operationally, Phase 4 reuses the Phase 3 baseline configuration `baseline_v1` and the floor ε recorded in `phase3/outputs/tables/mech_baseline_scan_diagnostics.json`. We then evaluate $A(\theta)$ and $E_{\mathrm{vac}}(\theta)$ on a uniform grid of $N_\theta = 2048$ points in $[0, 2\pi)$. The per-grid values and summary diagnostics are written to

`phase4/outputs/tables/phase4_F1_sanity_curve.csv`,

together with metadata describing the mapping parameters and the underlying Phase 3 diagnostics.

At this stage F1 is a *non-binding* mapping family: it does not yet define a θ-corridor or a Phase 4 θ-filter. Instead it serves as a concrete, auditable bridge between the Phase 3 mechanism and simple scalar observables that later rungs can connect to FRW-like toy modules and corridor construction.

3 Diagnostics and toy corridors

Given a mapping family from the Phase 3 global amplitude to a scalar observable, Phase 4 must provide diagnostics that are:

- simple enough to be implemented and audited end-to-end;
- honest about their physical status (toy vs. realistic); and
- compatible with the Phase 0 corridor / filter semantics.

At the present rung we focus on the F1 family, which maps the floor-enforced global amplitude $A(\theta)$ from Phase 3 into a toy vacuum-energy-like scalar

$$E_{\mathrm{vac}}(\theta) \;=\; \alpha\, A(\theta)^{\beta}, \tag{2}$$

with $\alpha = 1$ and $\beta = 4$ in the baseline configuration. We reuse the Phase 3 baseline diagnostics (including the quantile-based non-cancellation floor) and evaluate $E_{\mathrm{vac}}(\theta)$ on a uniform grid of $N_\theta = 2048$ points in $[0, 2\pi)$.

The per-grid values and a simple shape analysis are generated by the script

`phase4/src/phase4/run_f1_shape_diagnostics.py`,

which writes:

- a summary diagnostics file `phase4/outputs/tables/phase4_F1_shape_diagnostics.json`;
- a per-theta mask `phase4/outputs/tables/phase4_F1_shape_mask.csv`.

The diagnostics include:

- the global minimum and maximum of $E_{\mathrm{vac}}(\theta)$;

- the mean and standard deviation over the grid;
- a *toy, non-binding corridor* defined by

$$E_{\mathrm{vac}}(\theta) \ \le \ E_{\mathrm{vac,min}} \ + \ \sigma_E, \tag{3}$$

  where $E_{\mathrm{vac,min}}$ is the global minimum and σ_E is the standard deviation;
- the fraction of grid points inside this toy corridor; and
- the corresponding θ-interval spanned by the corridor.

This construction is explicitly labelled as a *non-binding* diagnostic: it is not a claim that the resulting interval defines a physically meaningful θ-corridor, let alone a canonical $\theta_\star$. Its purpose is to:

- demonstrate that the F1 mapping yields a numerically sensible θ-dependence;
- provide a reproducible starting point for more structured corridor definitions in later rungs; and
- illustrate how shape-based criteria on $E_{\mathrm{vac}}(\theta)$ can be turned into filters compatible with the Phase 0 ledger semantics, once the physical interpretation is better understood.

Later rungs will either refine these diagnostics into more physically motivated corridor conditions, or record structured negative results if no robust, non-pathological corridors emerge from the tested mapping families.

4 Limitations and outlook

Phase 4 is intentionally narrow in scope. Even once the mappings and diagnostics are implemented, the phase will not claim:

- a full derivation of cosmological parameters;
- a proof that the Origin Axiom is realised in nature; or
- a unique mechanism for connecting vacuum structure to FRW dynamics.

Instead, the goal is to provide a clean yes-or-no style test for a specific question:

> Can the Phase 3 global-amplitude mechanism support scale-sane FRW-like behaviour, in at least one simple mapping family, without producing a degenerate or empty θ-corridor?

If the answer is “no” for all tested mapping families, Phase 4 will record this as a structured negative result, signalling that either the Phase 3 mechanism or the mapping strategy needs revision before further unification attempts.

Appendix A: Phase 4 claims table (draft)

Table 1 summarises the intended Phase 4 claims. At this rung all entries are draft and non-binding.

Table 1: Draft Phase 4 claims. Binding status will be updated once the phase is complete and audited.

ID	Binding?	Summary
C4.1	no	Existence of at least one explicit mapping from the Phase 3 global amplitude or residue into an FRW-like or vacuum-energy-like observable with numerically stable behaviour.
C4.2	no	Existence of a non-empty, non-trivial θ-corridor for at least one such mapping.
C4.3	no	Structured negative result if all tested mappings yield empty or pathological corridors.

Appendix B: Reproducibility notes (draft)

This appendix will eventually document:

- the Phase 4 directory and workflow structure;
- the gate levels (paper-only vs. paper+artifacts);
- the commands required to rebuild the Phase 4 paper and any binding θ-filters; and
- the run manifests and configuration files that define the tested mapping families.

At this rung, Phase 4 only provides a minimal paper skeleton and no mapping implementations. Reproducibility therefore reduces to rebuilding the present PDF via the Phase 4 gate script.

References